

## RELATÓRIO ATIVIDADE 4 - INTELIGÊNCIA COMPUTACIONAL

### 1 Questão 1)

#### 1.1 Itens a) e b)

O código que implementa a construção, treinamento com gradiente descendente, e predição de uma rede neural de L camadas, com ativações  $\tanh()$ , está implementado no Algoritmo 1. Além disso, o código plota a evolução da função de custo MSE, assim como a evolução dos valores dos pesos e vies para cada neurônio. Também plota a região de classificação e a região de regressão. O Algoritmo 1 está no final do relatório.

### 2 Questão 2)

Para este exercício, utilizou-se a estrutura programada para a Questão 1, para programar a rede neural 2-2-1 e treiná-la para realizar a operação XOR. A parte que implementa esta questão e faz a plotagem dos gráficos está no final do Algoritmo 1.

#### 2.1 Item a)

As regiões de classificação e de regressão determinadas pelo algoritmo treinado estão representadas, respectivamente, nas Figuras 1 e 2. A região de classificação determina o pertencimento de cada ponto do espaço às classes 0 ou 1, sendo estas as saídas lógicas do XOR. Já a região de regressão retrata a gradação de valores de probabilidade de cada ponto. Esta probabilidade é levada em conta na hora de realizar a classificação final.

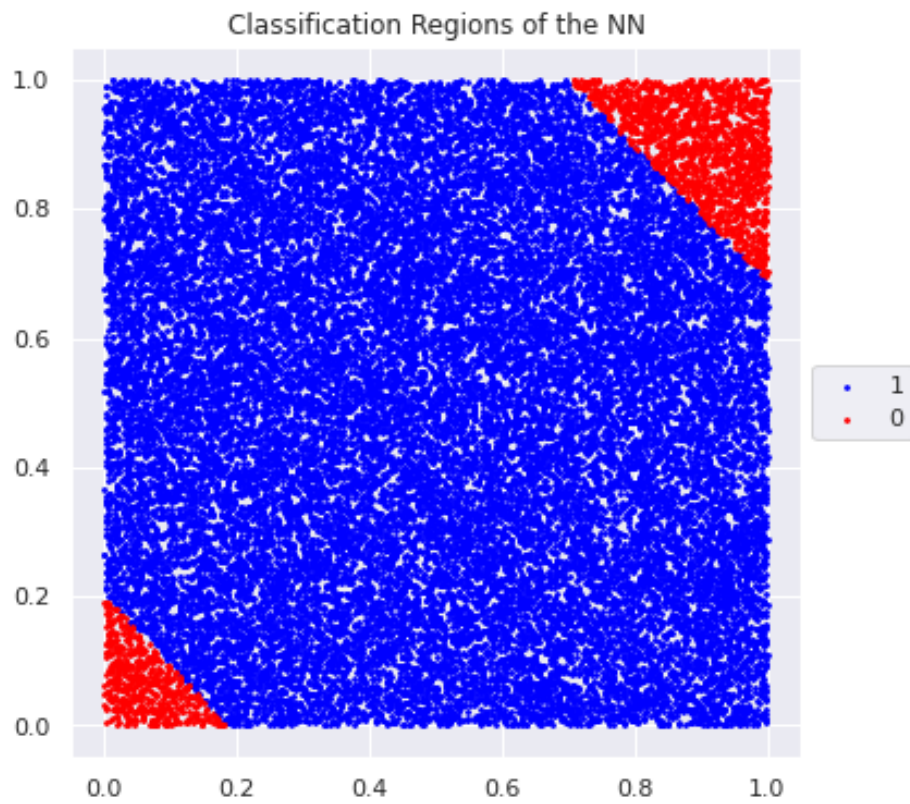


Figura 1: Regiões de classificação determinadas pelo algoritmo treinado.

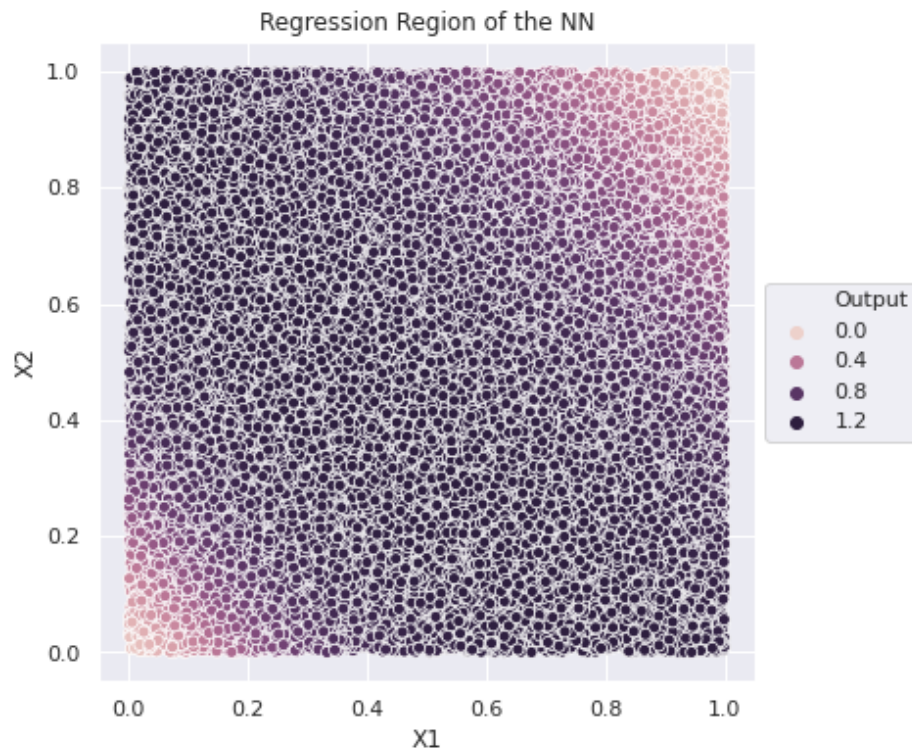


Figura 2: Regiões de regressão determinadas pelo algoritmo treinado.

## 2.2 Item b)

A evolução da função de custo MSE com as épocas de iteração está representada na Figura 3. Note que o custo atinge a tolerância de 0.001 em 372 iterações.




Figura 3: Evolução da função de custo MSE com as iterações de aprendizagem.

### 2.3 Item c)

Nas Figuras 4, 5 e 6 está exposto a evolução dos pesos e do viés para cada neurônio da rede, em relação ao número de iterações de aprendizagem do algoritmo.

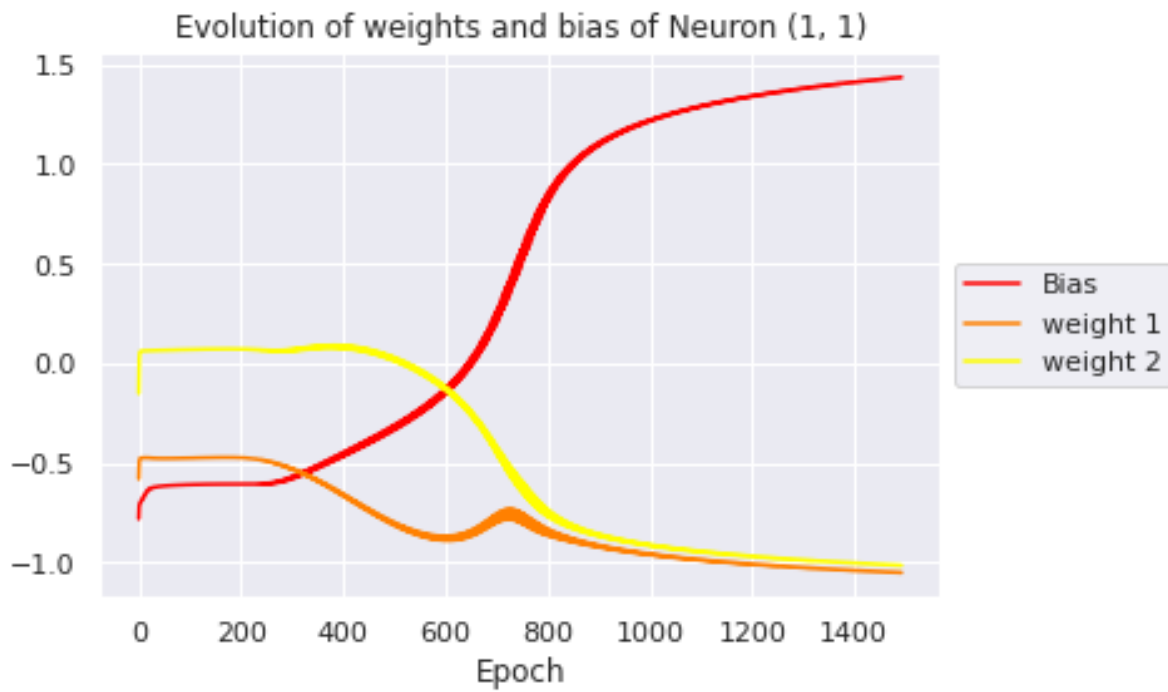


Figura 4: Evolução dos pesos e do viés para o primeiro neurônio da camada 1.

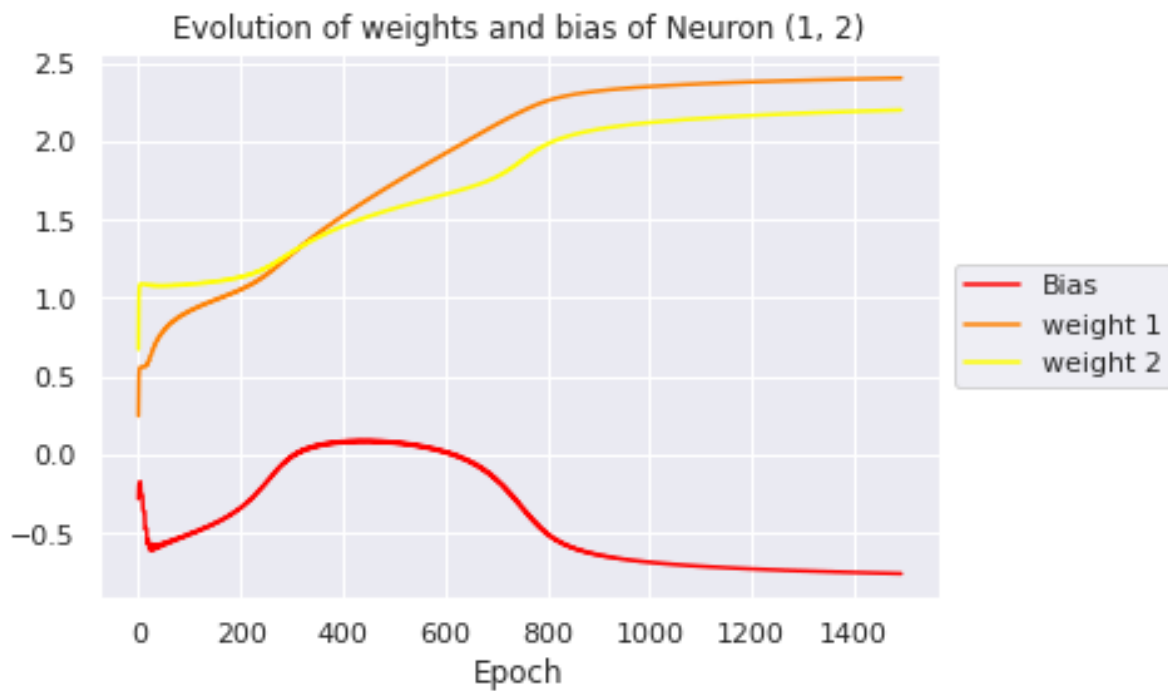


Figura 5: Evolução dos pesos e do viés para o segundo neurônio da camada 1.




Figura 6: Evolução dos pesos e do viés para o neurônio da camada 2 (neurônio de saída).

### 3 Algoritmo

```
# Andres E. M. Colognesi - 9838161

# #####
# ## Atividade 4 - Int. Comp. ##
# #####

# ----- Importing libraries -----

import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

sns.set()

# ----- Exercise 1 - Helper functions -----

def initialize_parameters_deep(layer_dims):
    """
```

```

Arguments:
layer_dims -- python array (list) containing the dimensions of
             each layer in our network

Returns:
parameters -- python dictionary containing your parameters
    "W1", "b1", ..., "WL", "bL":
        Wl -- weight matrix of shape (layer_dims[l],
            layer_dims[l-1])
        bl -- bias vector of shape (layer_dims[l], 1)
"""

np.random.seed(3)
parameters = {}
L = len(layer_dims)    # number of layers in the network

for l in range(1, L):
    parameters["W" + str(l)] = np.random.randn(layer_dims[l],
        layer_dims[l-1]) * 0.01
    parameters["b" + str(l)] = np.zeros((layer_dims[l], 1))
    #Verify dimensions:
    assert(parameters['W' + str(l)].shape == (layer_dims[l],
        layer_dims[l-1]))
    assert(parameters['b' + str(l)].shape == (layer_dims[l],
        1))

return parameters

def linear_forward(A, W, b):
    """
    Implement the linear part of a layer's forward propagation.

    Arguments:
    A -- activations from previous layer (or input data): (size of
        previous layer, number of examples)
    W -- weights matrix: numpy array of shape (size of current
        layer, size of previous layer)
    b -- bias vector, numpy array of shape (size of the current
        layer, 1)

    Returns:
    Z -- the input of the activation function, also called
        pre-activation parameter

```

```

cache -- a python dictionary containing "A", "W" and "b" ;
        stored for computing the backward pass efficiently
"""

Z = np.dot(W, A) + b
#Verify shape:
assert(Z.shape == (W.shape[0], A.shape[1]))
cache = (A, W, b)

return Z, cache

def my_tanh(Z):
    """
    Implements the tanh activation in numpy

    Arguments:
    Z -- numpy array of any shape

    Returns:
    A -- output of tanhh(z), same shape as Z
    cache -- returns Z as well, useful during backpropagation
    """

    A = np.tanh(Z)
    cache = Z

    return A, cache

def linear_activation_forward(A_prev, W, b, activation):
    """
    Implement the forward propagation for the LINEAR->ACTIVATION
    layer

    Arguments:
    A_prev -- activations from previous layer (or input data):
              (size of previous layer, number of examples)
    W -- weights matrix: numpy array of shape (size of current
              layer, size of previous layer)
    b -- bias vector, numpy array of shape (size of the current
              layer, 1)
    activation -- the activation to be used in this layer, stored
                  as a text string: "tanh"

```



```

Returns:
A -- the output of the activation function, also called the
    post-activation value
cache -- a python dictionary containing "linear_cache" and
    "activation_cache";
        stored for computing the backward pass efficiently
"""

if activation == "tanh":
    Z, linear_cache = linear_forward(A_prev, W, b)
    A, activation_cache = my_tanh(Z)
else:
    print("activation string error!")
#Verify dimensions:
assert (A.shape == (W.shape[0], A_prev.shape[1]))
cache = (linear_cache, activation_cache)

return A, cache

def L_model_forward(X, parameters):
    """
    Implement forward propagation for the
    [LINEAR->RELU]*(L-1)->LINEAR->SIGMOID computation

    Arguments:
    X -- data, numpy array of shape (input size, number of
        examples)
    parameters -- output of initialize_parameters_deep()

    Returns:
    AL -- last post-activation value
    caches -- list of caches containing:
        every cache of linear_activation_forward() (there
        are L-1 of them, indexed from 0 to L-1)
    """

    caches = []
    A = X
    L = len(parameters) // 2    # number of layers in the neural
        network

```

```

# Implement [LINEAR -> TANH]*(L). Add "cache" to the "caches"
list.
for l in range(1, L):
    A_prev = A
    A, cache = linear_activation_forward(A_prev,
        parameters["W"+ str(l)],parameters["b"+ str(l)],
        activation = "tanh")
    caches.append(cache)

# Implement LINEAR -> TANH. Add "cache" to the "caches" list.
AL, cache = linear_activation_forward(A, parameters["W"+
    str(L)],parameters["b"+ str(L)], activation = "tanh")
caches.append(cache)

assert(AL.shape == (1,X.shape[1]))

return AL, caches

def compute_cost(AL, Y):
    """
    Implement the cost function as MSE

    Arguments:
    AL -- probability vector corresponding to your label
        predictions, shape (1, number of examples)
    Y -- true "label" vector (for example: containing 0 if NO, 1
        if YES), shape (1, number of examples)

    Returns:
    cost -- cross-entropy cost
    """

    # Compute loss from aL and y.
    cost = np.mean( (Y-AL)**2 ) #MSE cost

    cost = np.squeeze(cost)      # To make sure your cost's shape
        is what we expect (e.g. this turns [[17]] into 17).
    assert(cost.shape == ())

    return cost

def linear_backward(dZ, cache):

```

```

"""
Implement the linear portion of backward propagation for a
    single layer (layer l)

Arguments:
dZ -- Gradient of the cost with respect to the linear output
    (of current layer l)
cache -- tuple of values (A_prev, W, b) coming from the
    forward propagation in the current layer

Returns:
dA_prev -- Gradient of the cost with respect to the activation
    (of the previous layer l-1), same shape as A_prev
dW -- Gradient of the cost with respect to W (current layer
    l), same shape as W
db -- Gradient of the cost with respect to b (current layer
    l), same shape as b
"""
A_prev, W, b = cache
m = A_prev.shape[1]

dW = 1/m*np.dot(dZ,A_prev.T)
db = 1/m*np.sum(dZ,axis=1,keepdims=True)
dA_prev = np.dot(W.T,dZ)
#Verify dimensions:
assert (dA_prev.shape == A_prev.shape)
assert (dW.shape == W.shape)
assert (db.shape == b.shape)

return dA_prev, dW, db

def tanh_backward(dA, cache):
    """
    Implement the backward propagation for a single TANH unit.
    Arguments:
    dA -- post-activation gradient, of any shape
    cache -- 'Z' where we store for computing backward propagation
        efficiently
    Returns:
    dZ -- Gradient of the cost with respect to Z
    """

    Z = cache

```

```

s, _ = my_tanh(Z)
dZ = dA * (1-s**2)
#Verify dimensions:
assert (dZ.shape == Z.shape)

return dZ

def linear_activation_backward(dA, cache, activation):
    """
    Implement the backward propagation for the LINEAR->ACTIVATION
    layer.

    Arguments:
    dA -- post-activation gradient for current layer l
    cache -- tuple of values (linear_cache, activation_cache) we
            store for computing backward propagation efficiently
    activation -- the activation to be used in this layer, stored
                 as a text string: "tanh"

    Returns:
    dA_prev -- Gradient of the cost with respect to the activation
               (of the previous layer l-1), same shape as A_prev
    dW -- Gradient of the cost with respect to W (current layer
         l), same shape as W
    db -- Gradient of the cost with respect to b (current layer
         l), same shape as b
    """
    linear_cache, activation_cache = cache

    if activation == "tanh":
        dZ = tanh_backward(dA, activation_cache)
        dA_prev, dW, db = linear_backward(dZ, linear_cache)
    else:
        print("error with backward activation function label!")

    return dA_prev, dW, db

def L_model_backward(AL, Y, caches):
    """
    Implement the backward propagation for the [LINEAR->TANH] *
    (L) group

```

Arguments:

AL -- probability vector, output of the forward propagation  
    (L\_model\_forward())  
Y -- true "label" vector (containing 0 if NO, 1 if YES)  
caches -- list of caches containing:  
          every cache of linear\_activation\_forward() with  
          "tanh"

Returns:

```
grads -- A dictionary with the gradients
        grads["dA" + str(l)] = ...
        grads["dW" + str(l)] = ...
        grads["db" + str(l)] = ...
"""
grads = {}
L = len(caches) # the number of layers
m = AL.shape[1]
Y = Y.reshape(AL.shape) # after this line, Y is the same shape
                           as AL

# Initializing the backpropagation
### START CODE HERE ### (1 line of code)
#dAL = - (np.divide(Y, AL) - np.divide(1 - Y, 1 - AL)) #
        derivative of cost with respect to AL
dAL = Y-AL
### END CODE HERE ###

# Lth layer (TANH -> LINEAR) gradients. Inputs: "dAL,
        current_cache". Outputs: "grads["dAL-1"], grads["dWL"],
        grads["dbL"]
current_cache = caches[L-1]
grads["dA" + str(L-1)], grads["dW" + str(L)], grads["db" +
        str(L)] = linear_activation_backward(dAL, current_cache,
        activation="tanh")

# Loop from l=L-2 to l=0
for l in reversed(range(L-1)):
    # lth layer: (TANH -> LINEAR) gradients.
    # Inputs: "grads["dA" + str(l + 1)], current_cache".
    # Outputs: "grads["dA" + str(l)] , grads["dW" + str(l +
        1)] , grads["db" + str(l + 1)]
    current_cache = caches[l]
```

```

    dA_prev_temp, dW_temp, db_temp =
        linear_activation_backward(grads["dA" + str(l + 1)],
            current_cache, activation="tanh")
    grads["dA" + str(l)] = dA_prev_temp
    grads["dW" + str(l + 1)] = dW_temp
    grads["db" + str(l + 1)] = db_temp

return grads

def update_parameters(parameters, grads, learning_rate,
    momentum_rate):
    """
    Update parameters using gradient descent

    Arguments:
    parameters -- python dictionary containing your parameters
    grads -- python dictionary containing your gradients, output
        of L_model_backward

    Returns:
    parameters -- python dictionary containing your updated
        parameters
        parameters["W" + str(l)] = ...
        parameters["b" + str(l)] = ...
    """

    L = len(parameters) // 2 # number of layers in the neural
        network

    # Update rule for each parameter. Use a for loop.
    for l in range(L):
        parameters["W" + str(l+1)] = parameters["W" + str(l+1)] -
            learning_rate * grads["dW" + str(l + 1)] -
            momentum_rate * (learning_rate * grads["dW" + str(l)])
        parameters["b" + str(l+1)] = parameters["b" + str(l+1)] -
            learning_rate * grads["db" + str(l + 1)] -
            momentum_rate * (learning_rate * grads["db" + str(l)])
    return parameters

def predict(X, y, parameters):
    """

```

This function is used to predict the results of a L-layer neural network.

Arguments:

X -- data set of examples you would like to label  
parameters -- parameters of the trained model

Returns:

p -- predictions for the given dataset X  
"""

```
m = X.shape[1]
n = len(parameters) // 2 # number of layers in the neural
    network
p = np.zeros((1,m))

# Forward propagation
probas, caches = L_model_forward(X, parameters)

# convert probas to 0/1 predictions
for i in range(0, probas.shape[1]):
    if probas[0,i] > 0:
        p[0,i] = 1
    else:
        p[0,i] = 0

#print results
print ("predictions: " + str(p))
print ("true labels: " + str(y))
print("Accuracy: " + str(np.sum((p == y)/m)))

return p
```

```
def L_layer_model(X, Y, layers_dims, learning_rate = 0.1,
    momentum_rate = 0.1, num_iterations = 3000, print_cost=False,
    tol=0.01):#lr was 0.009
    """
```

Implements a L-layer neural network:  
[LINEAR->RELU]\*(L-1)->LINEAR->SIGMOID.

Arguments:

```

X -- data, numpy array of shape (number of examples, entry
    size)
Y -- true "label" vector (containing 0 if NO, 1 if YES), of
    shape (1, number of examples)
layers_dims -- list containing the input size and each layer
    size, of length (number of layers + 1).
learning_rate -- learning rate of the gradient descent update
    rule
momentum_rate -- momentum rate of the gradient descent
    momentum update rule
num_iterations -- number of iterations of the optimization loop
print_cost -- if True, it prints the cost every step
tol -- accepted tolerance for error

Returns:
parameters -- parameters learnt by the model. They can then be
    used to predict.
"""

np.random.seed(1)
costs = []                                # keep track of cost

# Parameters initialization.
parameters = initialize_parameters_deep(layers_dims)

# Loop (gradient descent)
epoch = 0 #number of epochs
go_on = True #boolean to stop learning
while epoch < num_iterations and go_on:

    # Forward propagation: [LINEAR -> TANH]*(L).
    AL, caches = L_model_forward(X, parameters)

    # Compute cost.
    cost = compute_cost(AL, Y)

    # Backward propagation.
    grads = L_model_backward(AL, Y, caches)

    # Update parameters.
    parameters = update_parameters(parameters, grads,
        learning_rate, momentum_rate)

    # Print the cost every training example

```



```

    if print_cost:
        print ("Cost after iteration %i: %f" %(epoch, cost))
        costs.append(cost)

    if cost < tol: #the algorithm converged
        go_on = False #end learning loop
    epoch += 1
# plot the cost
plt.plot(np.squeeze(costs))
plt.xlabel("Epoch")
plt.ylabel("Mean Square Error")
plt.title("Learning curve of NN")
plt.show()

return parameters

def plotW(layer, pos, cmap='YlOrRd'):
    """
    Plots evolution of weights and bias.

    Arguments:
    layer : int - indicates neuron layer
    pos : int - indicates neuron position
    cmap : string (default = 'YlOrRd') - Colormap for the
           graphics.
    """

    w_array = np.array(self.w_list)
    n_iter, d = w_array.shape
    x = np.arange(n_iter)

    colors = eval(f"plt.cm.{cmap}(np.linspace(0,1,n))")

    names = ["Bias"]
    names += list(map(lambda i: f"Neuron {i}", np.arange(1,d)))

    for i in range(d):
        plt.plot(x, w_array[:,i], c=colors[i], label=names[i])
    plt.title(f"Evolution of weights and bias of Neuron
               ({layer, pos})")
    plt.legend(loc='center left', bbox_to_anchor=(1, 0.5))
    plt.xlabel("Epoch")
    plt.show()

```

```

#               ----- Exercise 2 - XOR -----

train_x = np.array([[0,0],
                    [0,1],
                    [1,0],
                    [1,1]]).T
train_y = np.array([[0,1,1,0]])

alpha = 0.1    # learning rate
mom = 0.1      # momentum rate
layers_dims = [2, 2, 1]    # 2 - layer model

parameters = L_layer_model(train_x, train_y, layers_dims,
                            learning_rate = alpha, momentum_rate = mom, num_iterations =
                            2500, print_cost = True)

test_data = np.random.random((30000,2))
pred_train = predict(train_x, train_y, parameters)

test_data0 = test_data[pred_train == 0]
test_data1 = test_data[pred_train == 1]

plt.figure(figsize=(6,6))
plt.scatter(test_data1[:,0], test_data1[:,1], s=3, c='red',
            label='1')
plt.scatter(test_data0[:,0], test_data0[:,1], s=3, c='blue',
            label='0')
plt.legend(loc='center left', bbox_to_anchor=(1, 0.5))
plt.title("Classification Regions of the NN")
plt.show()

df = pd.DataFrame(np.concatenate([test_data, predict(train_x,
            train_y, parameters)], axis=1), columns=['X1', 'X2', 'Output'])

plt.figure(figsize=(6,6))
sns.scatterplot(x='X1', y='X2', hue='Output', pallete = 'bright',
            data=df)
plt.legend(loc='center left', bbox_to_anchor=(1, 0.5))
plt.title("Regression Region of the NN")
plt.show()

```

```
# Plotting wheights and bias evolutions for each neuron:

for i in range(len(layers_dims)):
    for j in range(layers_dims[i]):
        plotW(i, j, 'autumn')
```

Algoritmo 1: Código em Python que implementa o algoritmo de foward e back propagation para uma rede de L camadas, com função de ativação *tanh* e função de custo MSE.